

Statistics

Assignment 2

1 Problem 3: Hypothesis testing

```
In [1]: import numpy as np
import pandas as pd
from math import sqrt
import scipy.stats as stats
import statsmodels.api as sm
import matplotlib.pyplot as plt
import matplotlib as mpl
import seaborn as sns
sns.set_style("darkgrid")
import warnings
warnings.filterwarnings('ignore')
```

1.1 1

We start with the verification of the law of large numbers. Thus we check if an estimator converges (in probability) to its true value if the sample size increases.

1.1.1 1.a

Simulate samples of size $n = 100, \dots, 100000$ from a normal distribution with mean 1 and variance 1, i.e. $N(1, 1)$. For each sample estimate the mean, the variance and store them.

```
In [138]: np.random.seed(0)

lognmax = 5
n_min = 10**2
n_max = 10**lognmax
ns = np.arange(n_min, n_max, 200)
num_experiments = len(ns)
print('{} samples in total. Some of them:'.format(num_experiments))
mu1 = 1
sigma1 = 1
var1 = sigma1**2
rv_norm1 = stats.norm(loc=mu1, scale=sigma1)
sample_means = []
sample_vars = []
```

```

prinit = 20
for i, n in enumerate(ns):
    sample = rv_norm1.rvs(size=n)
    sample_mean = sample.mean()
    sample_var = sample.var(ddof=1)
    sample_means.append(sample_mean)
    sample_vars.append(sample_var)
    if i % prinit == 0:
        print('N({}, {}) | Generated sample of size {:<{}}'\
              .format(mu1, sigma1**2, n, lognmax+1) + \
              ' | Est. Mean: {:.3f} | Est. Variance: {:.3f}'.format(sample_mean,
                                                                    sample_var))

```

500 samples in total. Some of them:

N(1,1)	Generated sample of size 100	Est. Mean: 1.060	Est. Variance: 1.026
N(1,1)	Generated sample of size 4100	Est. Mean: 1.004	Est. Variance: 0.964
N(1,1)	Generated sample of size 8100	Est. Mean: 1.011	Est. Variance: 0.997
N(1,1)	Generated sample of size 12100	Est. Mean: 1.011	Est. Variance: 0.984
N(1,1)	Generated sample of size 16100	Est. Mean: 0.997	Est. Variance: 1.009
N(1,1)	Generated sample of size 20100	Est. Mean: 0.992	Est. Variance: 0.997
N(1,1)	Generated sample of size 24100	Est. Mean: 1.000	Est. Variance: 0.989
N(1,1)	Generated sample of size 28100	Est. Mean: 1.001	Est. Variance: 1.002
N(1,1)	Generated sample of size 32100	Est. Mean: 0.997	Est. Variance: 0.999
N(1,1)	Generated sample of size 36100	Est. Mean: 0.997	Est. Variance: 0.995
N(1,1)	Generated sample of size 40100	Est. Mean: 1.003	Est. Variance: 1.004
N(1,1)	Generated sample of size 44100	Est. Mean: 0.995	Est. Variance: 0.992
N(1,1)	Generated sample of size 48100	Est. Mean: 1.001	Est. Variance: 0.992
N(1,1)	Generated sample of size 52100	Est. Mean: 1.001	Est. Variance: 1.007
N(1,1)	Generated sample of size 56100	Est. Mean: 0.999	Est. Variance: 1.001
N(1,1)	Generated sample of size 60100	Est. Mean: 0.995	Est. Variance: 0.990
N(1,1)	Generated sample of size 64100	Est. Mean: 1.001	Est. Variance: 0.999
N(1,1)	Generated sample of size 68100	Est. Mean: 0.999	Est. Variance: 0.994
N(1,1)	Generated sample of size 72100	Est. Mean: 0.996	Est. Variance: 1.000
N(1,1)	Generated sample of size 76100	Est. Mean: 0.998	Est. Variance: 1.008
N(1,1)	Generated sample of size 80100	Est. Mean: 0.998	Est. Variance: 1.005
N(1,1)	Generated sample of size 84100	Est. Mean: 0.998	Est. Variance: 1.003
N(1,1)	Generated sample of size 88100	Est. Mean: 0.995	Est. Variance: 0.989
N(1,1)	Generated sample of size 92100	Est. Mean: 1.003	Est. Variance: 1.002
N(1,1)	Generated sample of size 96100	Est. Mean: 0.992	Est. Variance: 1.001

Plot the path of sample means and sample variances as function of n.

```

In [167]: delta = 0.01
          mpl.rcParams['figure.figsize'] = (17,6)

          def plot_parameter(ns, estimated_values, true_value,
                              delta=None,

```

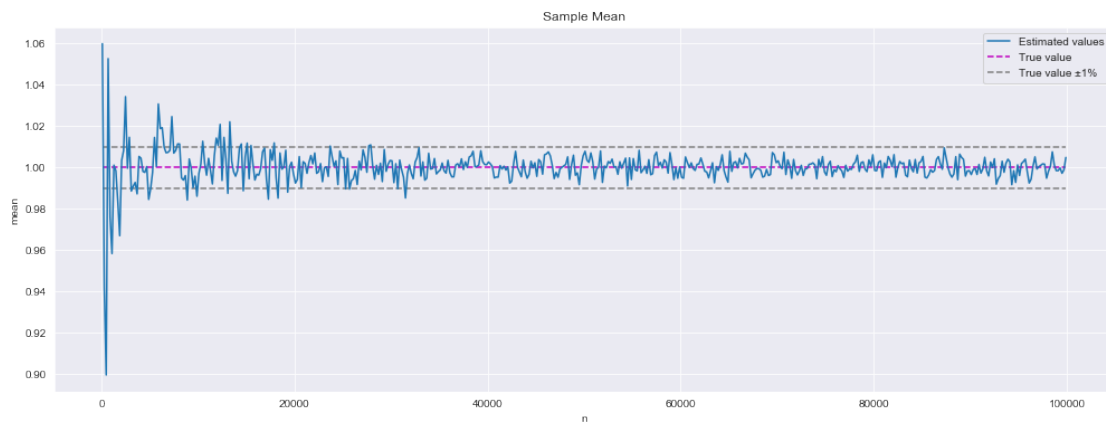
```


        cs=None, alpha=None,
        parameter_name=''):
    '''
    delta -- constant small deviation from the true_value;
            the horizontal lines with values of
            `(1 + delta) * true_value` will be plotted
    cs -- deviations to construct the confidence intervals
    '''

    plt.title('Sample {}'.format(parameter_name))
    plt.xlabel('n')
    plt.ylabel(parameter_name.lower())
    plt.plot(ns, estimated_values, label='Estimated values')
    plt.hlines(true_value, ns[0], ns[-1], 'm', '--', label='True value')
    if delta is not None:
        plt.hlines((1+delta)*true_value, ns[0], ns[-1], 'gray', '--',
                    label='True value  $\hat{s}{}%$ '.format(int(delta*100)))
        plt.hlines((1-delta)*true_value, ns[0], ns[-1], 'gray', '--')
    if cs is not None:
        if isinstance(cs, tuple):
            cs_lw, cs_up = cs[0], cs[1]
        else:
            cs_lw, cs_up = estimated_values - cs, estimated_values + cs
        plt.plot(ns, cs_up, '--', color='gray',
                  label='{}% confidence interval for the {}'\
                      .format(int((1-alpha)*100), parameter_name.lower()))
        plt.plot(ns, cs_lw, '--', color='gray')
    plt.legend()
    plt.show()

plot_parameter(ns, sample_means, mu1, delta=delta, parameter_name='Mean')
plot_parameter(ns, sample_vars, var1, delta=delta, parameter_name='Variance')

```





The plots support the LLN. As n becomes bigger the estimated values become closer to the true values.

How many observations do we need in order to obtain an estimator which is close enough (1%) to the true value?

```
In [173]: def min_num_observations(ns, estimated_values, true_value, delta):
            index = [(np.array(estimated_values)[i:] > (1-delta)*true_value).all() and
                      (np.array(estimated_values)[i:] < (1+delta)*true_value).all()
                      for i in range(len(ns))].index(True)
            return ns[index]

min_observations_mean = min_num_observations(ns, sample_means, mu1, delta)
min_observations_var = min_num_observations(ns, sample_vars, var1, delta)
min_observations = max(min_observations_mean, min_observations_var)
print('Experimentally we get that:')
print('\tAfter {:5} observations mean estimator is close enough (1%) to the true mean')
print('\tAfter {:5} observations variance estimator is close enough (1%) to the true variance')
print('\t=> After {:5} observations estimations are close enough to the true mean and variance')
```

Experimentally we get that:

After 31700 observations mean estimator is close enough (1%) to the true mean

After 95300 observations variance estimator is close enough (1%) to the true variance

=> After 95300 observations estimations are close enough to the true mean and variance

To get the number of observations theoretically we should use the known true value of σ and set the confidence level to $1 - \alpha$ e.g. 0.95:

$$P\left(|\mu - \bar{X}| < z_{1-\alpha/2} \frac{\sigma}{\sqrt{n}}\right) > 1 - \alpha \quad (1)$$

We want $|\mu - \bar{X}| < 0.01\mu = 0.01$

$$\Rightarrow z_{1-\alpha/2} \frac{\sigma}{\sqrt{n}} = 0.01 \quad (2)$$

$$n = z_{1-\alpha/2}^2 \frac{\sigma^2}{0.0001} \quad (3)$$

```
In [180]: alpha = 0.05
          n = stats.norm.ppf(1 - alpha/2) ** 2 * var1 / (0.01*mu1)**2
          print('We need more than {:.0f} observations to obtain an estimator close enough (±1%
```

We need more than 38415 observations to obtain an estimator close enough (±1%) to the true value

1.b Add to the plot the 95% confidence intervals. These have to be constructed manually. Provide their interpretation.

For the mean the confidence interval is provided with unknown σ (more often we don't know the true variance):

$$\left(\bar{X} - t_{n-1;1-\alpha/2} \frac{S}{\sqrt{n}}; \bar{X} + t_{n-1;1-\alpha/2} \frac{S}{\sqrt{n}} \right) \quad (4)$$


Analogously the confidence interval for the variance is constructed as we don't know the true mean:

$$\left(\frac{(n-1)S^2}{\chi_{n-1;1-\alpha/2}^2}; \frac{(n-1)S^2}{\chi_{n-1;\alpha/2}^2} \right) \quad (5)$$

```
In [172]: alpha = 0.05
```

```
cs = np.array([stats.t.ppf(1-alpha/2, df=n-1) * np.sqrt(S2/n)
               for (n,S2) in zip(ns, sample_vars)])
plot_parameter(ns, sample_means, mu1, cs=cs, alpha=alpha,
               parameter_name='Mean')

cs_lw = np.array([S2*(n-1) / stats.chi2.ppf(1-alpha/2, df=n-1)
                  for (n,S2) in zip(ns, sample_vars)])
cs_up = np.array([S2*(n-1) / stats.chi2.ppf(alpha/2, df=n-1)
                  for (n,S2) in zip(ns, sample_vars)])
plot_parameter(ns, sample_vars, var1, cs=(cs_lw, cs_up),
               alpha=alpha, parameter_name='Variance')
```



With 95% confidence we are sure that the true parameter value is inside the confidence interval.

1.2 2

The 2nd objective of this part is to get more feeling for the ML estimation procedures. The estimation for non-standard distributions/models usually follows the maximum likelihood principle. The t-distribution is a popular alternative if the sample distribution is symmetric but exhibits heavier tails compared to the normal distribution.

2 (a) Let x_1, \dots, x_n be a given sample. We assume that it stems from a t-distribution with an unknown number of degrees of freedom. Write down the corresponding likelihood function.

$$X_1, \dots, X_n \sim t_d : \quad f(x_i) = \frac{(1 + \frac{x_i^2}{d})^{-\frac{d+1}{2}}}{B(d/2, 1/2)\sqrt{d}} \quad (6)$$

$$L_d(x_1, \dots, x_n) = \prod_{i=1}^n \frac{(1 + \frac{x_i^2}{d})^{-\frac{d+1}{2}}}{B(d/2, 1/2)\sqrt{d}} = \frac{\prod_{i=1}^n (1 + \frac{x_i^2}{d})^{-\frac{d+1}{2}}}{B^n(d/2, 1/2) \cdot d^{n/2}} \quad (7)$$

$$\ln L_d(x_1, \dots, x_n) = \frac{-\frac{d+1}{2} \sum_{i=1}^n \ln(1 + \frac{x_i^2}{d})}{n \ln(B(d/2, 1/2) \cdot \sqrt{d})} \quad (8)$$

2 (b) Simulate a sample of size $n = 100$ from t_5 . Maximize the likelihood function (numerically) for the given sample and obtain the ML estimator of the number of degrees of freedom.

```
In [18]: def log_likelihood(pdfs):
         return np.sum(np.log(pdfs))

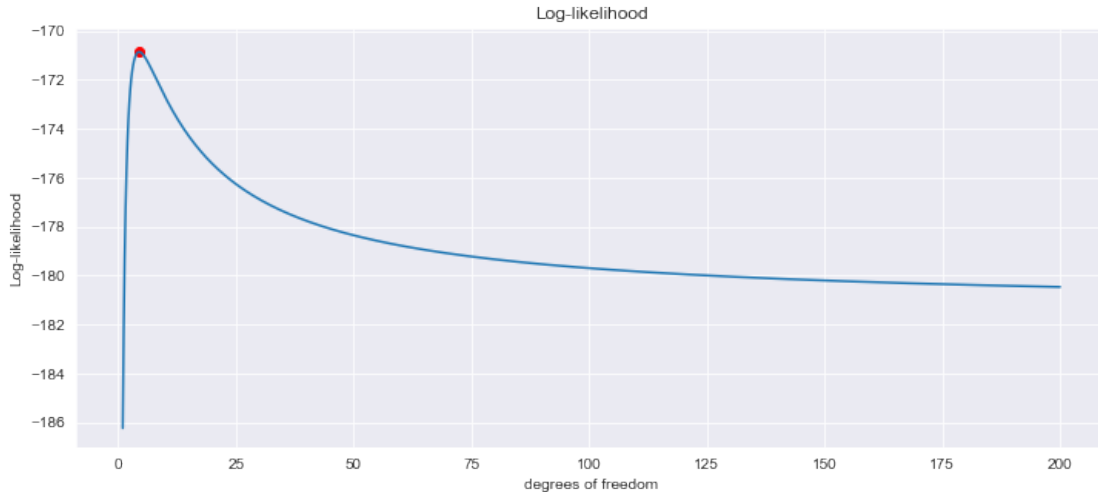
In [185]: mpl.rcParams['figure.figsize'] = (12, 5)
         np.random.seed(100)

         d_true = 5
         n = 100
         rv_st = stats.t(df=d_true)
         sample_st = rv_st.rvs(size=n)

         ds = np.linspace(1,200,1000)
         Ls = []
         for d in ds:
             pdfs = stats.t(df=d).pdf(sample_st)
             # plt.scatter(sample_st, pdfs)
             Ls.append(log_likelihood(pdfs))
         # plt.show()
         d_est = ds[np.argmax(Ls)]
         print('df* = {:.2f} maximizes likelihood function'.format(d_est))

         plt.title('Log-likelihood')
         plt.xlabel('degrees of freedom')
         plt.ylabel('Log-likelihood')
         plt.plot(ds, Ls)
         plt.scatter(d_est, np.max(Ls), color='red')
         plt.show()

         df* = 4.59 maximizes likelihood function
```



2 (c) The classical theoretical t-distribution has zero mean and variance $df/(df-2)$. A real sample does not have exactly these moments. How would you proceed if you need to fit a t-distribution to real data?

If we transform random variable X s.t. $Y = \sigma X + \mu$ then $F_Y(x) = F_X(\frac{x-\mu}{\sigma})$ and $f_Y(x) = F'_X(\frac{x-\mu}{\sigma}) = \frac{1}{\sigma} f_X(\frac{x-\mu}{\sigma})$

We can use it with transformed Student r.v. to calculate likelihood:

$$L(x_1, \dots, x_n) = \prod_{i=1}^n \frac{1}{\sigma} f_X\left(\frac{x_i - \mu}{\sigma}\right) \quad (9)$$

As was shown in previous assignment we can transform r. v. X from Student distribution with n degrees of freedom s.t. it has mean μ and variance σ^2 : $Y = \sigma \sqrt{\frac{n-2}{n}} X + \mu$. So we can estimate n using ML and thus fit t_n -distribution to our data.

```
In [197]: mpl.rcParams['figure.figsize'] = (12, 5)
          np.random.seed(100)

          d_true = 5
          n = 100

          mu = 10
          sigma = 5
          sample_st = mu + sigma * np.sqrt((df-2) / df) * stats.t(df=d_true).rvs(size=n)

          ds = np.linspace(1,200,1000)
          Ls = []
          for d in ds:
              pdfs = stats.t(df=d).pdf((sample_st - mu) / sigma) / sigma
```



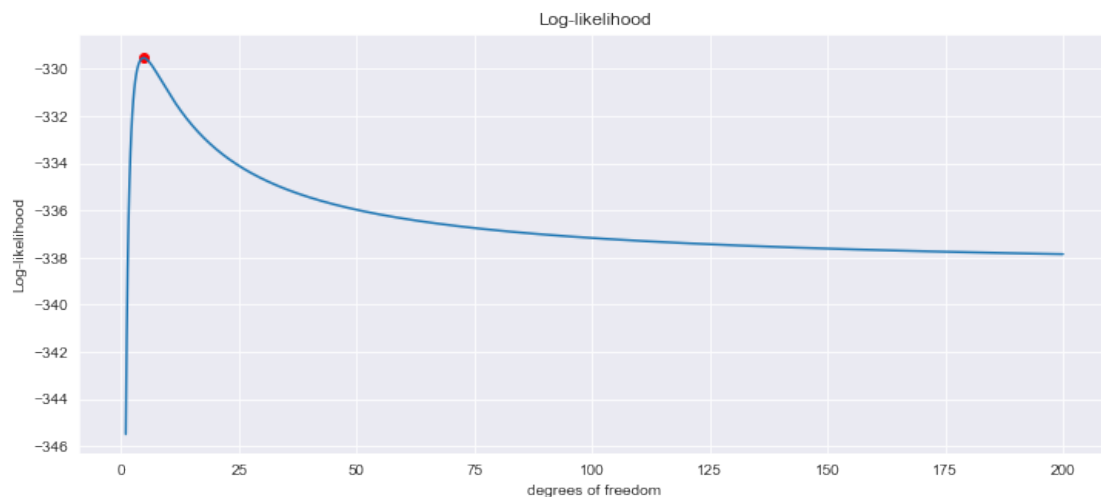
```

    Ls.append(log_likelihood(pdfs))
d_est = ds[np.argmax(Ls)]
print('df* = {:.2f} maximizes likelihood function'.format(d_est))

plt.title('Log-likelihood')
plt.xlabel('degrees of freedom')
plt.ylabel('Log-likelihood')
plt.plot(ds, Ls)
plt.scatter(d_est, np.max(Ls), color='red')
plt.show()

```

df* = 4.78 maximizes likelihood function



1.3 3

The 3rd objective is check if the probability of type 1 error (size of a test) is correctly attained by a simple two-sided test for the mean.

3 (a) Simulate a sample of length $n = 100$ from a normal distribution with mean $\mu_0 = 500$ and variance $\sigma_2 = 50$. (Note: you may use the transformation $X = \mu + \sigma Z$, where $Z \sim N(0,1)$.) The objective is to test the null hypothesis $H_0 : \mu = 500$. Assume that σ_2 has to be estimated. Compute the test statistics using the formulas in the lecture; determine the rejection area for $\alpha = 0.04$ and decide if H_0 can to be rejected.

In [23]: `np.random.seed(0)`

```

n = 100
mu = 500
sigma = 50

```

```
rv = stats.norm(loc=mu, scale=sigma)
x = rv.rvs(size=n)
```

```
In [24]: def normal_cdf(x):
    return stats.norm().cdf(x)

def normal_ppf(F):
    return stats.norm().ppf(F)

def student_cdf(df, x):
    return stats.t(df=df).cdf(x)

def student_ppf(df, F):
    return stats.t(df=df).ppf(F)

def two_sided_Z_test(x, a, alpha=None, sigma=None, out=True):
    '''
    H0: mu == a
    H1: mu != a
    alpha -- significance level
    '''
    if out:
        print('\nTwo-sided Z-test {}'.format('(sigma is unknown)'
                                              if sigma is None else ''))

        print('H0: mu == {}'.format(a))
        print('H1: mu != {}'.format(a))
    n = len(x)
    sample_mean = np.mean(x)
    # print('\tSample mean: {:.2f}'.format(sample_mean))
    if sigma is None:
        sample_var = np.var(x, ddof=1)
        # print('\tSample variance: {:.2f}'.format(sample_var))
        v = (sample_mean - a) * sqrt(n) / sqrt(sample_var)
        p_value = 2 * student_cdf(n-1, v) if v < 0 else \
            2 * (1 - student_cdf(n-1, v))
        if alpha is not None:
            t_crit = student_ppf(n-1, 1-alpha/2)
    else:
        v = (sample_mean - a) * sqrt(n) / sigma
        p_value = 2 * normal_cdf(v) if v < 0 else 2 * (1 - normal_cdf(v))
        if alpha is not None:
            t_crit = normal_ppf(1-alpha/2)

    if alpha is not None:
        rejected = v < -t_crit or v > t_crit
    if out:
        print('The rejection area is: (-inf, -{:.3f}) ({:.3f}, inf)\'\\
```

```

        .format(t_crit, t_crit))
    print('t_stat = {:.3f} is {} in the rejection area.'\
          .format(v, ' if rejected else 'not'))
    if not rejected:
        print('With significance level of {}% we cannot reject H0.'\
              .format(alpha*100, a))
    else:
        print('With significance level of {}% we can reject H0.')
    return p_value, rejected
if out:
    print('p-value is {:.3f}'.format(p_value))
    print('For all > {:.3f} we can reject H0.'.format(p_value))
return p_value

_, _ = two_sided_Z_test(x, mu, 0.04)

```

Two-sided Z-test (sigma is unknown)

H0: $\mu == 500$

H1: $\mu != 500$

The rejection area is: $(-\infty, -2.081) \quad (2.081, \infty)$

t_stat = 0.590 is not in the rejection area.

With significance level of 4.0% we cannot reject H0.

3 (b) Determine the p-values using the formulas from the lecture and compare/check the results using a build-in function for this test in R or Python.

In [31]: `p_value = two_sided_Z_test(x, mu)`

```

tstat, pvalue = stats.ttest_1samp(x, mu)
print('\np-value from `scipy.stats` is {:.3f}'.format(pvalue))

```

Two-sided Z-test (sigma is unknown)

H0: $\mu == 500$

H1: $\mu != 500$

p-value is 0.556

For all > 0.556 we can reject H0.

p-value from `scipy.stats` is 0.556

So implemented p-value is the same as within `scipy.stats`.

3 (c) Simulate $M = 1000$ samples of size $n = 100$ and with $\mu_0 = 500$ and variance $\sigma^2 = 50$. For each sample i run the test (using a standard function) and set $p_i = 0$ if H_0 is not rejected and $p_i = 1$ if rejected. Compute $\hat{\alpha} = \frac{1}{M} \sum_{i=1}^M p_i$. $\hat{\alpha}$ is the empirical condence level (empirical size) of

the test. Compare $\hat{\alpha}$ to α ? Do you expect the difference to be large or small and why? Relate it to the assumptions of the test.

```
In [111]: M = 1000
          n = 100
          mu = 500
          sigma = np.sqrt(50)
          alpha = 0.04

          rv = stats.norm(loc=mu, scale=sigma)
          X = rv.rvs(size=n*M).reshape((M,n))
          P = []
          for x in X:
              _, pvalue = sm.stats.ztest(x, value=mu)
              rejected = alpha > pvalue
              P.append(int(rejected))

In [112]: print('Empirical condence level: {}, level of significance {}'.format(np.mean(P), alpha))
```

Empirical condence level: 0.044, level of significance 0.04

The expectations are that the empirical confidence level would be close to α . Since we assume that data is normally distributed we should get the type 1 error with probability α . So for our 1000 experiments we should get roughly $1000\alpha = 40$ rejections, indeed we got 44. The bigger is M the closer $\hat{\alpha}$ should be to α , due to the Law of Large Numbers.

3 (d) Assume now that one of the assumptions is not satisfied. For example, the data is in fact not normal. Repeat the above analysis, but simulate the data from the t-distribution with 3 degrees of freedom. Give motivation and justification for the new values of $\hat{\alpha}$?

```
In [123]: M = 1000
          n = 100
          df = 3

          rv = stats.t(df=df)
          X = rv.rvs(size=n*M).reshape((M,n))
          X = mu + sigma * np.sqrt((df-2) / df) * X
          P = []
          for x in X:
              _, pvalue = sm.stats.ztest(x, value=mu)
              rejected = alpha > pvalue
              P.append(int(rejected))

In [124]: print('Empirical condence level: {}'.format(np.mean(P), alpha))
```

Empirical condence level: 0.033

We should also get empirical value relatively close to the level of significance since the distribution of the data is close to normal. Nevertheless the assumptions on the normal distribution are not fulfilled so we may need a bigger number of experiments to get a very close estimate.

3 (e) Power of a test: The first objective is to assess the probability of type 2 error (power of a test) of goodness-of-t test. Goodness-of-t tests for the normal distribution are of key importance in statistics, since they allow to verify the distributional assumptions required in many models. Here we check the power of the Kolmogorov-Smirnov test, i.e. is the test capable to detect deviations from normality?

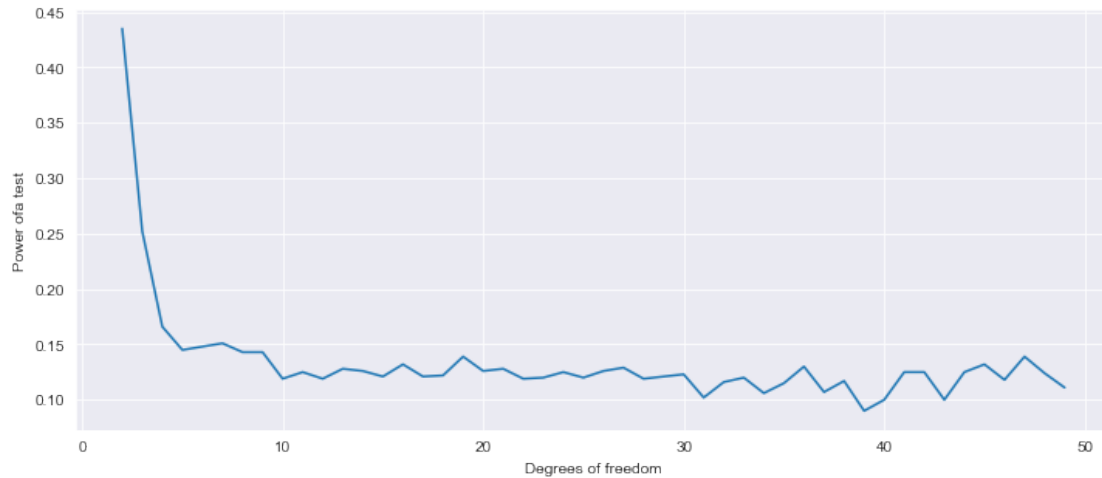
Simulate $M = 1000$ samples of size 100 from a t-distribution with $df = 2, \dots, 50$ degrees of freedom. For each sample run the Kolmogorov-Smirnov test and count the cases when the H_0 of normality is correctly rejected (for each df). How would you use this quantity to estimate the power of the test? Make an appropriate plot with the df on the X-axis. (Note: the t-distribution converges to the normal distribution as df tends to infinity. For $df > 50$ the distributions are almost identical. Discuss the plot and draw conclusions about the reliability of the test.

```
In [131]: np.random.seed(0)
```

```
M = 1000
n = 100
dfs = np.arange(2,50)
print('H0: X ~ Normal')
print('H1: X !~ Normal')
alphas = []
for df in dfs:
    rv = stats.t(df=df)
    X = rv.rvs(size=n*M).reshape((M,n))
    P = []
    for x in X:
        ksstat, pvalue = stats.kstest(x, 'norm')
        rejected = ksstat > pvalue
        P.append(int(rejected))
    alphas.append(np.mean(P))
```

```
H0: X ~ Normal
H1: X !~ Normal
```

```
In [135]: plt.plot(dfs, alphas)
plt.ylabel('Power of a test')
plt.xlabel('Degrees of freedom')
plt.show()
```



The data is not distributed normally so H_0 of the Kolmogorov-Smirnov test should be rejected. A test's power ($1 - \beta$) is the probability that the test rejects the H_0 when H_1 is true. So the frequency of rejection H_0 with our data can be used as a direct estimation of a test's power. From the plot we see that as df grows the H_0 is rejected less frequently, meaning that the test's power decreases. This is expectable since the t-distribution becomes closer to the normal distribution with increasing df.

Problem4

1 Problem 4: Linear regression analysis

A telephone service provider aims to decrease the churn rate and analyses the data and service usage of 1000 clients. The following variables are used in the study

- tenure - month a client;
- age - age in years;
- marital status - marital status (1 - married, 0 - single);
- address - years at the current address;
- income - household income in Euro;
- ed - education (5 categories: Did not complete high school; High school degree; Some college; College degree; Post-undergraduate degree);
- retire - retired (0 - no, 1 - yes);
- gender - gender (0 - male, 1 - female);
- longmon - long distance calls last month;
- wiremon - internet use last month;
- churn - 1 if the contract was terminated last month and 0 else

The overall objective is to analyze the service usage using longmon as the dependent variables and the remaining variables as explanatory.

```
In [489]: import numpy as np
import pandas as pd
import scipy as sp
import scipy.stats as stats
import statsmodels.api as sm
from sklearn import linear_model, metrics
import matplotlib.pyplot as plt
import matplotlib as mpl
import seaborn as sns
sns.set_style("darkgrid")
import warnings
warnings.filterwarnings('ignore')
from utils import *

df_raw = pd.read_csv('data/telco.txt', sep='\t')
N = len(df_raw)
df_raw.head()
```

```

Out[489]:
   tenure  age  marital  address  income  ed \
1      13   44   Married      9      64      College degree
2      11   33   Married      7     136      Post-undergraduate degree
3      68   52   Married     24     116      Did not complete high school
4      33   33  Unmarried     12      33      High school degree
5      23   30   Married      9      30      Did not complete high school

   employ  retire  gender  longmon  wiremon  churn
1         5      No    Male     3.70      0.0    Yes
2         5      No    Male     4.40     35.7    Yes
3        29      No  Female    18.15      0.0     No
4         0      No  Female     9.45      0.0    Yes
5         2      No    Male     6.30      0.0     No

```

1.1 1.

Have a closer look at the denitions of the variables and analyze which of them might require a separate treatment. Consider for example the variable `ed`. There are two possibilities how the variable `ed` can be included into the model (one with dummy variables, the other one without dummies). Think about these two approaches and suggest which approach is more appropriate. Motivate your decision.

```
In [2]: print('Data types:\n{}'.format(df_raw.dtypes))
```

Data types:

```

tenure      int64
age         int64
marital     object
address     int64
income      int64
ed          object
employ      int64
retire      object
gender      object
longmon     float64
wiremon     float64
churn       object
dtype: object

```

```

In [3]: # Group variables by the type of the scale
        intervals = list(df_raw.dtypes[(df_raw.dtypes == np.int64) |
                                         (df_raw.dtypes == np.float64)].index)
        # ['tenure', 'age', 'address', 'income',
        #  'employ', 'longmon', 'wiremon']

        ordinals = ['ed']

        nominals = list(set(df_raw.dtypes.index).difference(intervals, ordinals))

```



```

# ['marital', 'retire', 'gender', 'churn']

df_intervals = df_raw[intervals]
df_ordinals = df_raw[ordinals]
df_nominals = df_raw[nominals]

In [4]: # Look at interval scaled variables
stats_descr = sp.stats.describe(df_intervals)
describe_intervals = df_intervals.describe(percentiles=[.1, .25, .5, .75, .9]).append(
    [pd.Series(stats_descr.skewness, index=intervals, name='skew'),
     pd.Series(stats_descr.variance, index=intervals, name='var')])
describe_intervals = describe_intervals.drop('count',0)
describe_intervals

```

```

Out [4]:

```

	tenure	age	address	income	employ \
mean	35.526000	41.684000	11.551000	77.535000	10.987000
std	21.359812	12.558816	10.086681	107.044165	10.082087
min	1.000000	18.000000	0.000000	9.000000	0.000000
10%	7.000000	26.000000	1.000000	21.000000	0.000000
25%	17.000000	32.000000	3.000000	29.000000	3.000000
50%	34.000000	40.000000	9.000000	47.000000	8.000000
75%	54.000000	51.000000	18.000000	83.000000	17.000000
90%	66.000000	59.000000	26.100000	155.400000	25.000000
max	72.000000	77.000000	55.000000	1668.000000	47.000000
skew	0.111692	0.356128	1.104586	6.633303	1.059457
var	456.241566	157.723868	101.741140	11458.453228	101.648479

	longmon	wiremon
mean	11.723100	11.583900
std	10.363486	19.719426
min	0.900000	0.000000
10%	3.645000	0.000000
25%	5.200000	0.000000
50%	8.525000	0.000000
75%	14.412500	24.712500
90%	23.960000	42.110000
max	99.950000	111.950000
skew	2.961653	1.601274
var	107.401848	388.855747

```

In [5]: # Look at nominal scaled variables
df_nominals.describe()

```

```

Out [5]:

```

	marital	churn	retire	gender
count	1000	1000	1000	1000
unique	2	2	2	2
top	Unmarried	No	No	Female
freq	505	726	953	517

```
In [6]: # Look at ordinal scaled variables
df_ordinals.describe()
```

```
Out [6]:
```

	ed
count	1000
unique	5
top	High school degree
freq	287

The variable ed can be represented as dummy variables:

```
In [435]: ed_dummies = pd.get_dummies(df_raw.ed)
ed_names_dict = {'College degree': 'ed_college',
                 'Did not complete high school': 'ed_no',
                 'High school degree': 'ed_highschool',
                 'Post-undergraduate degree': 'ed_postgr',
                 'Some college': 'ed_somecollege'}
ed_dummies = ed_dummies.rename(columns=ed_names_dict).drop(['ed_no'], 1)
ed_dummies.head()
```

```
Out [435]:
```

	ed_college	ed_highschool	ed_postgr	ed_somecollege
1	1	0	0	0
2	0	0	1	0
3	0	0	0	0
4	0	1	0	0
5	0	0	0	0

```
In [8]: df_dummies = pd.concat((df_raw.drop('ed', 1), ed_dummies), 1)
df_dummies['marital'] = df_dummies.marital.map(dict(Married=1, Unmarried=0))
df_dummies['retire'] = df_dummies.retire.map(dict(Yes=1, No=0))
df_dummies['gender'] = df_dummies.gender.map(dict(Female=1, Male=0))
df_dummies['churn'] = df_dummies.churn.map(dict(Yes=1, No=0))
df_dummies.head()
```

```
Out [8]:
```

	tenure	age	marital	address	income	employ	retire	gender	longmon	\
1	13	44	1	9	64	5	0	0	3.70	
2	11	33	1	7	136	5	0	0	4.40	
3	68	52	1	24	116	29	0	1	18.15	
4	33	33	0	12	33	0	0	1	9.45	
5	23	30	1	9	30	2	0	0	6.30	

	wiremon	churn	ed_college	ed_highschool	ed_postgr	ed_somecollege
1	0.0	1	1	0	0	0
2	35.7	1	0	0	1	0
3	0.0	0	0	0	0	0
4	0.0	1	0	1	0	0
5	0.0	0	0	0	0	0


Another option is to use the property of the order scaled variable and map the values to their ranks:

```

In [457]: ed_order_dict = {'Did not complete high school': 0,
                           'High school degree': 1,
                           'Some college': 2,
                           'College degree': 3,
                           'Post-undergraduate degree': 4}

df_ordered = df_raw.copy()
df_ordered['ed'] = df_ordered.ed.map(ed_order_dict)
mpl.rcParams['figure.figsize'] = (15, 3)
plt.hist(df_ordered.ed, align='mid')
plt.xticks([0.2,1,2.2,3,3.8], ed_order_dict.keys())
plt.show()
df_ordered.head()

```



```

Out [457]:
   tenure  age  marital  address  income  ed  employ  retire  gender \
1      13   44   Married      9      64   3        5      No   Male
2      11   33   Married      7     136   4        5      No   Male
3      68   52   Married     24     116   0       29      No  Female
4      33   33  Unmarried     12      33   1        0      No  Female
5      23   30   Married      9      30   0        2      No   Male

   longmon  wiremon  churn
1      3.70      0.0   Yes
2      4.40     35.7   Yes
3     18.15      0.0   No
4      9.45      0.0   Yes
5      6.30      0.0   No

```

On the one hand, the second approach should be more appropriate since we include the information about the order of ed e.g. some education is better or higher than no education at all.

On the other hand with this order we "fix" the difference between degrees and say e.g. that Post-undergraduate degree is higher than College degree as much as Some college is higher than High school degree etc. while this may not be true. But with dummy variables this is not fixed, so it is suggested to use futher.

```

In [10]: df = df_dummies.copy()

```

1.2 2.

Consider now the dependent variable `longmon` and the interval (metric) scaled explanatory variables. Plot these data and decide if you wish to transform these x-variables and if there is a need to transform the y variable. You can also use some measure of skewness to decide about y. The variable `wiremon` shows a very specific pattern. How would you take it into account?

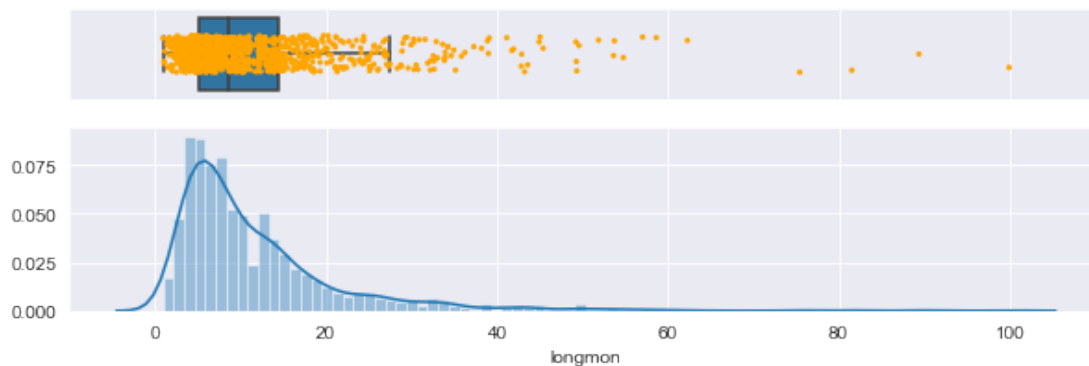
```
In [488]: mpl.rcParams['figure.figsize'] = (17,5)
```

longmon (long distance calls last month)

```
In [12]: column = 'longmon'
         represent_distribution(df[column], varname=column)
         print(describe_intervals.longmon)
         skewness = sp.stats.skew(df[column])
         print('The distribution is {}'.format('right-skewed' if skewness > 0 else 'left-skewed'))
```

Variable `longmon`

Number of observations: 1000



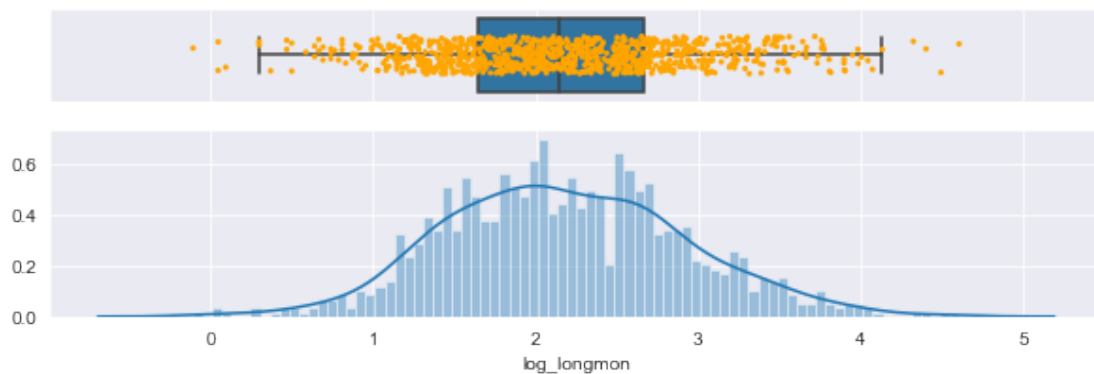
```
mean      11.723100
std       10.363486
min        0.900000
10%        3.645000
25%        5.200000
50%        8.525000
75%       14.412500
90%       23.960000
max       99.950000
skew       2.961653
var       107.401848
Name: longmon, dtype: float64
The distribution is right-skewed.
```

We transform this variable longmon into log(longmon):

```
In [13]: df['log_{}'.format(column)] = np.log(df[column])
         column = 'log_{}'.format(column)
         represent_distribution(df[column], varname=column)
         # stats_descr = sp.stats.describe(df[['log_longmon']])
         # df[['log_longmon']].describe(percentiles=[.1, .25, .5, .75, .9]).append(
         # [pd.Series(stats_descr.skewness, index=['log_longmon'], name='skew'),
         # pd.Series(stats_descr.variance, index=['log_longmon'], name='var')]).drop('count', 0)
```

Variable `log_longmon`

Number of observations: 1000

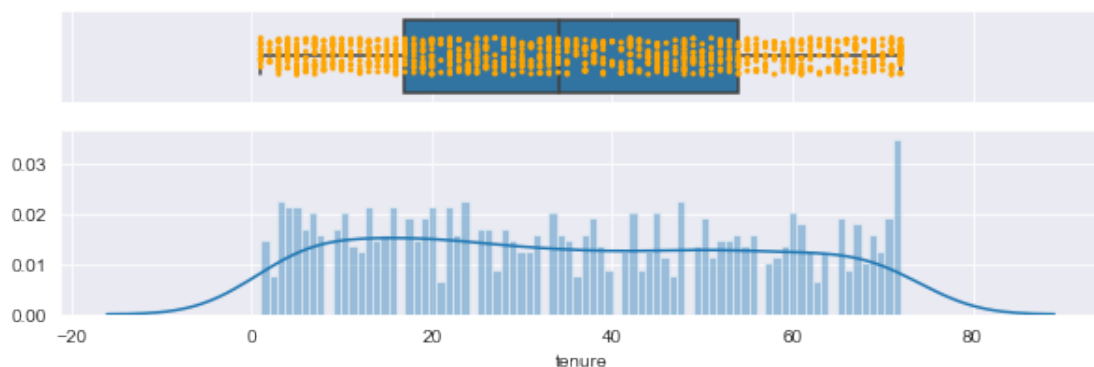


tenure

```
In [14]: column = 'tenure'
         represent_distribution(df[column], varname=column)
         print(describe_intervals[column])
```

Variable `tenure`

Number of observations: 1000



```

mean      35.526000
std       21.359812
min        1.000000
10%        7.000000
25%       17.000000
50%       34.000000
75%       54.000000
90%       66.000000
max       72.000000
skew      0.111692
var      456.241566
Name: tenure, dtype: float64

```

age (age in years)

```


In [15]: column = 'age'
         represent_distribution(df[column], varname=column)
         print(describe_intervalls[column])

```

```

Variable `age`
Number of observations: 1000

```



```

mean      41.684000
std       12.558816
min       18.000000
10%       26.000000
25%       32.000000
50%       40.000000
75%       51.000000

```

```

90%      59.000000
max       77.000000
skew      0.356128
var       157.723868
Name: age, dtype: float64

```

address (years at the current address)

```


In [16]: column = 'address'
         represent_distribution(df[column], varname=column)
         print(describe_intervals[column])

```

```

Variable `address`
Number of observations: 1000

```



```

mean      11.551000
std       10.086681
min        0.000000
10%        1.000000
25%        3.000000
50%        9.000000
75%       18.000000
90%       26.100000
max       55.000000
skew       1.104586
var       101.741140
Name: address, dtype: float64


```

income (household income in Euro)

```
In [17]: column = 'income'
         represent_distribution(df[column], varname=column)
         print(describe_intervals[column])
```

Variable `income`

Number of observations: 1000



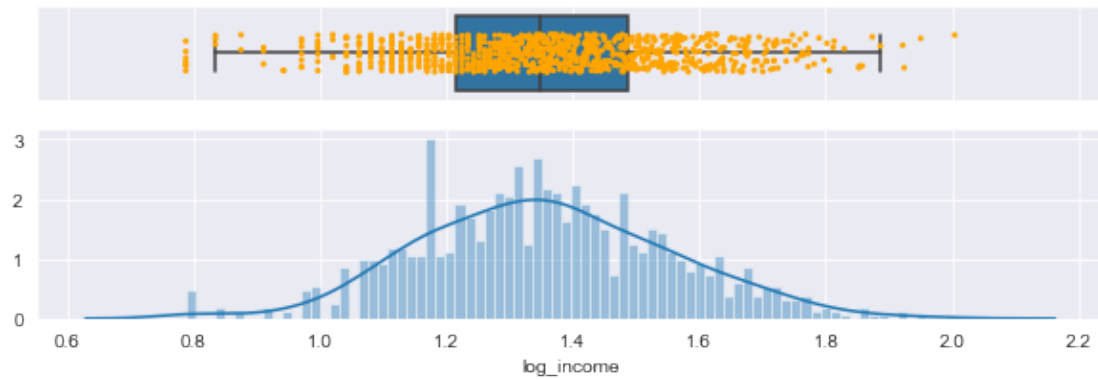
```
mean      77.535000
std      107.044165
min        9.000000
10%       21.000000
25%       29.000000
50%       47.000000
75%       83.000000
90%      155.400000
max      1668.000000
skew       6.633303
var      11458.453228
Name: income, dtype: float64
```

As with longmon we transform the variable income into $\log(\text{income})$:

```
In [18]: df['log_{}'.format(column)] = np.log(df[column])
         column = 'log_{}'.format(column)
         represent_distribution(np.log(df[column]), varname=column)
```

Variable `log_income`


Number of observations: 1000



employ (years with the current employer)

```
In [19]: column = 'employ'
         represent_distribution(df[column], varname=column)
         print(describe_intervals[column])
```

Variable `employ`
Number of observations: 1000




mean	10.987000
std	10.082087
min	0.000000
10%	0.000000
25%	3.000000
50%	8.000000
75%	17.000000
90%	25.000000
max	47.000000

```
skew      1.059457
var       101.648479
Name: employ, dtype: float64
```

wiremon (internet use last month)

```
In [20]: column = 'wiremon'
         represent_distribution(df[column], varname=column)
         print(describe_intervals[column])
```

```
Variable `wiremon`
Number of observations: 1000
```




```
mean      11.583900
std       19.719426
min        0.000000
10%        0.000000
25%        0.000000
50%        0.000000
75%       24.712500
90%       42.110000
max       111.950000
skew      1.601274
var       388.855747
Name: wiremon, dtype: float64
```

Let's look at the distribution of usage for those who used the internet last month:

```
In [21]: represent_distribution(df[df[column]>0][column], varname=column+' > 0')
```


```
Variable `wiremon > 0`
Number of observations: 296
```



We can create a dummy variable representing, whether the client used or not the internet last month:

```
In [22]: df['wire_not_used'] = (df['wiremon']==0).astype(np.int)
          nominals.append('wire_not_used')
```

```
In [487]: df[nominals].hist(figsize=(20, 10))
          plt.show()
```



```
In [24]: df.head()
```

```
Out[24]:
```

	tenure	age	marital	address	income	employ	retire	gender	longmon	\
1	13	44	1	9	64	5	0	0	3.70	
2	11	33	1	7	136	5	0	0	4.40	

3	68	52	1	24	116	29	0	1	18.15
4	33	33	0	12	33	0	0	1	9.45
5	23	30	1	9	30	2	0	0	6.30

	wiremon	churn	ed_college	ed_highschool	ed_postgr	ed_somcollege	\
1	0.0	1	1	0	0	0	
2	35.7	1	0	0	1	0	
3	0.0	0	0	0	0	0	
4	0.0	1	0	1	0	0	
5	0.0	0	0	0	0	0	

	log_longmon	log_income	wire_not_used
1	1.308333	4.158883	1
2	1.481605	4.912655	0
3	2.898671	4.753590	1
4	2.246015	3.496508	1
5	1.840550	3.401197	1

1.3 3.

After making up your decision about the above two problems run a simple linear regression.

```
In [25]: def R_squared(y, y_hat):
    y = np.array(y).reshape(-1,1)
    y_hat = np.array(y_hat).reshape(-1,1)
    y_mean = y.mean()
    ESS = np.linalg.norm(y_hat - y_mean)**2
    TSS = np.linalg.norm(y - y_mean)**2
    return ESS/TSS

def R_squared_adj(y, y_hat, K):
    y = np.array(y).reshape(-1,1)
    y_hat = np.array(y_hat).reshape(-1,1)
    N = len(y)
    y_mean = y.mean()
    RSS = np.linalg.norm(y - y_hat)**2
    TSS = np.linalg.norm(y - y_mean)**2
    return 1 - RSS * (N-1) / (TSS * (N-K-1))

def AIC(y, y_hat, K):
    y = np.array(y).reshape(-1,1)
    y_hat = np.array(y_hat).reshape(-1,1)
    u = y - y_hat
    s2 = np.var(u, ddof=1)
    N = len(y)
    return np.log(s2) + 2 * K / N

def BIC(y, y_hat, K):
```

```

y = np.array(y).reshape(-1,1)
y_hat = np.array(y_hat).reshape(-1,1)
u = y - y_hat
N = len(y)
s2 = np.var(u, ddof=1)
return np.log(s2) + K * np.log(N) / N

In [26]: def get_X_y(df, regressand, regressors, out=True):
    if out:
        print('Regressand:', regressand)
        print('Regressors:', regressors)
    y = df[regressand]
    X = df[regressors]
    return X, y

def get_normalized_X_y(X, y, intervals_regressors):
    X_intervals_means = X[intervals_regressors].mean()
    X_intervals_stds = X[intervals_regressors].std(ddof=1)

    X_norm = X.copy()
    X_norm[intervals_regressors] = (X[intervals_regressors] - X_intervals_means) / X_intervals_stds

    y_mean = y.mean()
    y_std = y.std(ddof=1)
    y_norm = (y - y_mean) / y_std

    return X_norm, y_norm, X_intervals_means, X_intervals_stds, y_mean, y_std

def get_LR_beta(X, y, out_scores=False):
    LR = linear_model.LinearRegression()
    LR.fit(X, y)
    beta_df = pd.DataFrame([LR.intercept_]+list(LR.coef_),
                           index=X.columns.insert(0, 1),
                           columns=['coef'])

    if out_scores:
        print('LR scores:')
        get_scores(LR, X, y)
    return LR, beta_df

def get_scores(LR, X, y):
    # y_hat = np.dot(sm.add_constant(X), beta_df['coef'].values.reshape(-1,1))
    y_hat = LR.predict(X)
    R2 = R_squared(y, y_hat)
    R2adj = R_squared_adj(y, y_hat, K)
    aic = AIC(y, y_hat, K)
    bic = BIC(y, y_hat, K)
    print('R^2 = {:.3f}\nR^2_adj = {:.3f}\nAIC = {:.3f}\nBIC = {:.3f}'.format(R2, R2adj, aic, bic))

```

```
def get_standardized_coef(beta_df_orig, intervals_regressors, other_regressors):
    # For self-check
    beta_df = beta_df_orig.copy()
    beta_df['coef_st'] = beta_df['coef']
    beta_df.loc[intervals_regressors, 'coef_st'] = beta_df.loc[intervals_regressors, 'coef'] / y
    beta_df.loc[other_regressors, 'coef_st'] = beta_df.loc[other_regressors, 'coef'] / y
    beta_df.loc[1, 'coef_st'] = (beta_df.loc[1, 'coef'] + np.sum(beta_df.loc[intervals_regressors, 'coef'])) / y
    return beta_df
```

```
In [707]: # Simple LR
regressand = 'log_longmon'
# with wiremon
# regressors = sorted(list(set(df.columns).difference({'longmon', 'log_longmon', 'income'})))
# intervals_regressors = sorted(list(set(intervals).union({'log_income'}).difference({'longmon', 'log_longmon', 'income'})))

# with wire_used
regressors = sorted(list(set(df.columns).difference({'longmon', 'log_longmon', 'income'})))
intervals_regressors = sorted(list(set(intervals).union({'log_income'}).difference({'longmon', 'log_longmon', 'income'})))
other_regressors = sorted(list(set(regressors).difference(intervals_regressors)))

X, y = get_X_y(df, regressand, regressors)
K = X.shape[1]
```

Regressand: log_longmon

Regressors: ['address', 'age', 'churn', 'ed_college', 'ed_highschool', 'ed_postgr', 'ed_somecol']

Before running LR we normalize X and y , and work only with it. As will be shown in Problem 5, it won't affect R^2 and we can get coefficients for the unnormalized data with corresponding linear transformations.

```
In [709]: X, y, X_intervals_means, X_intervals_stds, y_mean, y_std = \
            get_normalized_X_y(X, y, intervals_regressors)

X_ = sm.add_constant(X)

LR, beta_df = get_LR_beta(X, y)
beta_df['coef_abs'] = np.abs(beta_df['coef'])

# Get standardized coefficients of LR
# LR_st, beta_df = get_LR_beta(X_norm, y_norm, True)
# beta_df['coef_st'] = beta_st['coef']
# beta_df['coef_st_abs'] = np.abs(beta_df['coef_st'])
print('\nStandardized LR coefficients:')
print(beta_df.sort_values(by='coef_abs', ascending=False)[['coef']])
print()
y_hat = LR.predict(X)
print('R_squared: {:.3f}'.format(R_squared(y, y_hat)))
print('R_squared_adj: {:.3f}'.format(R_squared_adj(y, y_hat, K)))
print('AIC: {:.3f}'.format(AIC(y, y_hat, K)))
```

```
print('BIC: {:.3f}'.format(BIC(y, y_hat, K)))
```

```
LR = sm.OLS(y, X_)
LR_results = LR.fit()
# print(LR_results.summary())
```

Standardized LR coefficients:

	coef
tenure	0.793197
retire	0.334490
ed_college	0.157259
1	-0.126201
ed_highschool	0.124676
ed_somecollege	0.108959
marital	0.098148
churn	-0.054113
address	0.051514
gender	-0.048634
age	-0.047336
employ	0.024508
log_income	0.023402
ed_postgr	0.012869
wire_not_used	0.008047

R_squared: 0.717

R_squared_adj: 0.713

AIC: -1.235

BIC: -1.167

If you wish to argue that education is insignificant and use the model with dummies then you have to check the simultaneous insignificance of all dummies which stem from the factor variable ed. Run a test for general linear hypothesis and conclude about the significance of ed.

Variable ed seems to be important, let's check it:

```
In [31]: ed_vars = ['ed_somecollege', 'ed_college', 'ed_highschool', 'ed_postgr']
```

```
def test_significance(LR_results, variables):
    H0 = ' = '.join(variables) + ' = 0'
    if len(variables) > 1:
        H0_text = 'all in {} simultaneously have no impact'.format(variables)
        H1_text = 'at least one in {} is significant'.format(variables)
    else:
        H0_text = '{} has no impact'.format(variables)
        H1_text = '{} is significant'.format(variables)
    print('H0:', H0_text)
    print('H1:', H1_text)
```

```

F_results = LR_results.f_test(H0)
print('F = {:.4f}, p-value = {:.4f}'.format(F_results.fvalue[0,0], F_results.pvalue[0,0]))

test_significance(LR_results, ed_vars)

H0: all in ['ed_somcollege', 'ed_college', 'ed_highschool', 'ed_postgr'] simultaneously have no effect
H1: at least one in ['ed_somcollege', 'ed_college', 'ed_highschool', 'ed_postgr'] is significant
F = 2.7484, p-value = 0.0272

```

The p-value is small enough so we decide to reject H0, which means that the variable education is significant.

1.4 4.

Provide an economic interpretation for the parameters of address, ed, and retire. Neglect the possible insignificance and keep in mind possible transformations of the variables.

```

In [32]: beta_df.loc[['address', 'retire',
                    'ed_somcollege', 'ed_college',
                    'ed_highschool', 'ed_postgr'], :]\
        .sort_values(by='coef_abs', ascending=False)[['coef']]

```

```

Out[32]:
           coef
retire      0.334490
ed_college  0.157259
ed_highschool 0.124676
ed_somcollege 0.108959
address     0.051514
ed_postgr   0.012869

```

```

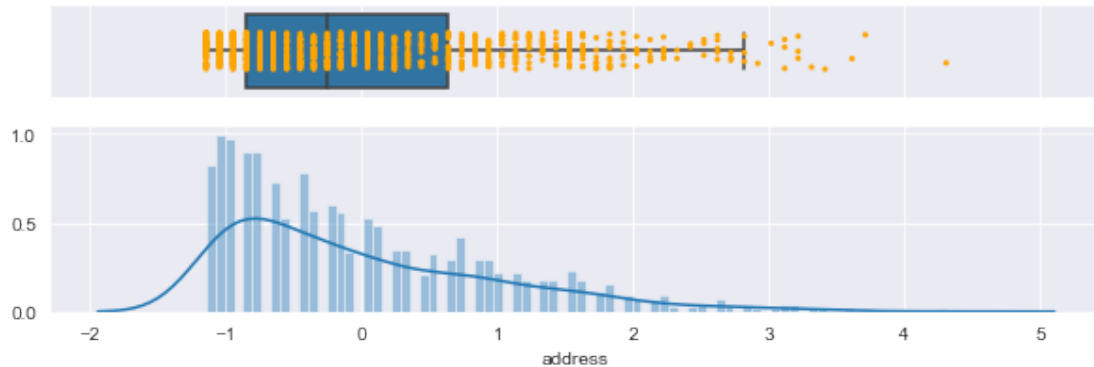
In [33]: print('Address mean: {:.2f}, std: {:.2f} (used for normalization)'.format(X_intervals['address'].mean(), X_intervals['address'].std()))
represent_distribution(X['address'], varname='normalized address')
test_significance(LR_results, ['address'])
print()
test_significance(LR_results, ['retire'])

```

```

Address mean: -0.00, std: 1.00 (used for normalization)
Variable `normalized address`
Number of observations: 1000

```

H0: ['address'] has no impact
H1: ['address'] is significant
F = 4.6722, p-value = 0.0309

H0: ['retire'] has no impact
H1: ['retire'] is significant
F = 9.5569, p-value = 0.0020

The p-value for the first test with address is small so this variable is significant. This variable represents years at the current address but during the normalization 11.5 was subtracted from values and it was divided by 10 so now it's hard to interpret this variable. The p-value for the test with retire is also very small so this variable is significant. It is a nominal scaled variable representing whether a person is retired or not, and it wasn't changed. The coefficient 0.3345 tells that with the same other parameters change from 'not retired' to 'retired' would make us expect an increasing in normalized log_longmon by 0.3345. For education we already concluded that it is a significant variable. It represents type of education and, as was supposed, we see that indeed a change on different education levels causes a different change in target, and also the order of (sorted) coefficients is not the same as a natural order of education levels (where e.g. college should be between highschool and postgr) so the relation is more complex.

1.5 5.

Compute the 95% condence intervals for the parameters of address and income and provide its economic meaning. Relate the CIs to the tests of signigance, i.e. how would you use these intervals to decide about the signigance of the corresponding explanatory variables? The CIs are computed relying on the assumption, that the residuals follow normal distribution. Is this assumption fulfilled? Run an appropriate goodness-of-fit test.

```
In [34]: print('95%-CI:')
         CI = LR_results.conf_int(alpha=0.05).loc[['address', 'log_income'], :]\
           .rename(columns={0: 'lower', 1: 'upper'})
```

```

CI['length'] = CI.upper - CI.lower
print(CI)
print()
test_significance(LR_results, ['log_income'])

```

95%-CI:

	lower	upper	length
address	0.004746	0.098283	0.093537
log_income	-0.033571	0.080375	0.113946

H0: ['log_income'] has no impact

H1: ['log_income'] is significant

F = 0.6497, p-value = 0.4204

The lower bound for the log_income is negative, the upper bound is positive. So zero value for the coefficient is inside 95% confidence interval which mean possible insignificance of the variable. The significance test results in quite large p-value 0.42, so we can reject the null-hypothesis (stating that the variable is significant) and drop log_income. We already concluded that address is significant. We see that the CI's length for this variable is less than for previous meaning it has less uncertainty.

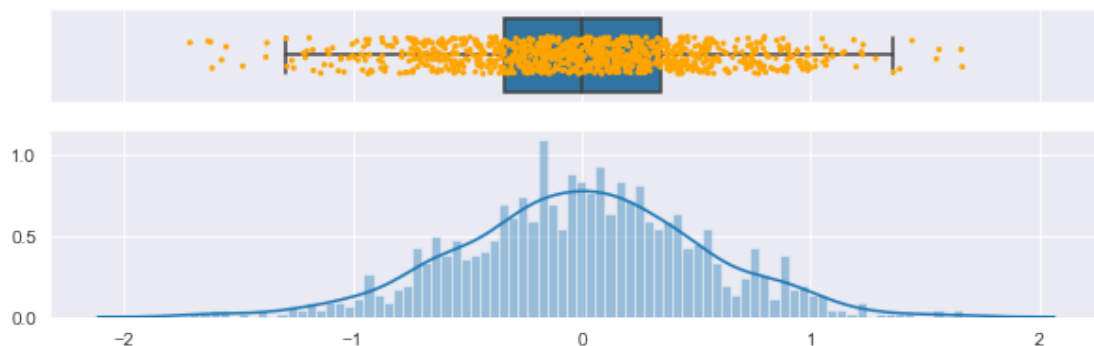
```

In [35]: print('Residuals')
          res = LR_results.resid
          represent_distribution(res)
          print('Mean: {:.2e}'.format(res.mean()))

```

Residuals

Number of observations: 1000



Mean: -3.03e-15

The histogram of the residuals looks like it follows the normal distribution with zero-mean. To check the assumption on the normal distribution we will run Kolmogorov-Smirnov goodness-of-fit test.

```
In [48]: print('Kolmogorov-Smirnov test for residuals.')
print('H0: res ~ Normal')
print('H1: res !~ Normal')
_, pvalue = stats.kstest(res/res.std(), 'norm')
print('p-value: {:.2f}'.format(pvalue))
```

```
Kolmogorov-Smirnov test for residuals.
H0: res ~ Normal
H1: res !~ Normal
p-value: 0.85
```

With the p-value 0.85 we do not reject the hypothesis of normal distributed residuals. Therefore computing CI is reasonable.

1.6 6.

Many of the variable appear insignificant and we should find the smallest model, which still has a good explanatory power. Choose this model using stepwise model selection (either based on the tests for R^2 or using AIC/BIC). Pick up the last step of the model selection procedure and explain in details how the method/approach works (or is implemented in your software). Work with this model in all the remaining steps.

```
In [147]: mpl.rcParams['figure.figsize'] = (20,5)
print('In total {} regressors: {}'.format(len(regressors), regressors))

def forward_selection(X, y, regressors, score_type, verbose=False, plot=True):
    score_fn = R_squared_adj if score_type == 'R^2_adj' else AIC if score_type == 'AIC' else BIC
    regressors_selected, regressors_left = [], regressors.copy()
    scores_selected = []
    best_score = np.inf if score_type in ('AIC', 'BIC') else -np.inf
    score = best_score
    while regressors_left != [] and score==best_score:
        regressor_score = []
        for regressor in regressors_left:
            X_selected = X.loc[:, regressors_selected+[regressor]]
            LR = linear_model.LinearRegression()
            LR.fit(X_selected, y)
            y_hat = LR.predict(X_selected)
            s = score_fn(y, y_hat, len(X_selected.columns))
            regressor_score.append(s)
        if verbose:
            print(regressors_selected+[regressor], ': {:.3f}'.format(s))
        i = np.argmin(regressor_score) if score_type in ('AIC', 'BIC') \
            else np.argmax(regressor_score)
        score = regressor_score[i]
        condition = (score < best_score) if score_type in ('AIC', 'BIC') \
            else (score > best_score)
        if condition:
```

```

        if verbose:
            print('> {} is chosen with score {}'.format(regressors_left[i], score))
            best_score = score
            regressors_selected.append(regressors_left[i])
            scores_selected.append(score)
            del regressors_left[i]
    print('> In terms of {} ({}):\n{}'.format(score_type, len(regressors_selected),
    if plot:
        xs = regressors_selected+regressors_left[i]
        xs = ['{} \\'{}\''.format('' if i==0 else '+', x) for i, x in enumerate(xs)]
        plt.plot(xs, scores_selected+[score], '--')
        plt.scatter(xs[:-1], scores_selected)
        plt.scatter(xs[-1], score, color='r', zorder=10)
        plt.title('{} score'.format(score_type))
        plt.show()
    return regressors_selected

print('\nThe best model (forward selection):')
forward_selection(X, y, regressors, 'R^2_adj')
regressors_selected = forward_selection(X, y, regressors, 'AIC')
forward_selection(X, y, regressors, 'BIC');


```

In total 14 regressors: ['address', 'age', 'churn', 'ed_college', 'ed_highschool', 'ed_postgr

The best model (forward selection):


> In terms of R²_adj (11):

['tenure', 'retire', 'marital', 'ed_college', 'ed_highschool', 'ed_somecollege', 'address', 'cl




> In terms of AIC (7):

['tenure', 'retire', 'marital', 'ed_college', 'ed_highschool', 'ed_somecollege', 'address']



> In terms of BIC (2):
['tenure', 'retire']



The result of forward model selection using AIC was chosen as a middle ground. The selected regressors are: tenure, retire, marital, ed_college, ed_highschool, ed_somecollege, address.

```
In [148]: intervals_regressors = list(set(regressors_selected).intersection(set(intervals_regr
other_regressors = list(set(regressors_selected).intersection(set(other_regressors)))
X = X.loc[:, regressors_selected]
X_ = sm.add_constant(X)
LR = sm.OLS(y, X_)
LR_results = LR.fit()
print(LR_results.summary())
```

OLS Regression Results

Dep. Variable:	log_longmon	R-squared:	0.715
Model:	OLS	Adj. R-squared:	0.713
Method:	Least Squares	F-statistic:	355.7
Date:	Tue, 26 Feb 2019	Prob (F-statistic):	2.36e-265
Time:	13:38:25	Log-Likelihood:	-790.60
No. Observations:	1000	AIC:	1597.

Df Residuals: 992 BIC: 1636.
Df Model: 7
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-0.1508	0.038	-3.969	0.000	-0.225	-0.076
tenure	0.8079	0.020	39.798	0.000	0.768	0.848
retire	0.2546	0.084	3.019	0.003	0.089	0.420
marital	0.0927	0.035	2.675	0.008	0.025	0.161
ed_college	0.1515	0.048	3.138	0.002	0.057	0.246
ed_highschool	0.1229	0.046	2.698	0.007	0.034	0.212
ed_somecollege	0.1061	0.050	2.136	0.033	0.009	0.204
address	0.0405	0.021	1.962	0.050	-4.99e-06	0.081
Omnibus:	2.592	Durbin-Watson:	1.944			
Prob(Omnibus):	0.274	Jarque-Bera (JB):	2.634			
Skew:	-0.054	Prob(JB):	0.268			
Kurtosis:	3.227	Cond. No.	6.24			

Warnings:



[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.7 7.

Sometimes data contains outliers which induces bias in the parameter estimates. Check for outliers using Cook's distance and leverage. Have a closer look at the observation with the highest leverage (regardless if it is classied as an outlier or not). What makes this observation so outstanding (you may have a look at Box-plots for interval scaled variables or at the frequencies for binary/ordinal variables?

```
In [207]: from statsmodels.stats.outliers_influence import OLSInfluence


influence = OLSInfluence(LR_results).summary_frame()
pvals = LR_results.get_influence().cooks_distance[1]
plt.stem(influence.standard_resid, influence.cooks_d, markerfmt=",")
plt.title('Cook\'s distance')
plt.show()
plt.stem(influence.standard_resid, influence.cooks_d, markerfmt=",")
plt.plot(influence.standard_resid, pvals, 'k--')
plt.title('Cook\'s distance (with p-values)')
plt.show()
```



```
In [216]: plt.title('Leverage')
leverage = influence.hat_diag
plt.stem(influence.standard_resid, leverage, markerfmt="," , basefmt='gray')

n = 50
ind = leverage.index[(-leverage).argsort()][:n]
plt.stem(influence.standard_resid.loc[ind], leverage.loc[ind],
        markerfmt="r," , linefmt='r', basefmt='gray')

plt.legend(['leverage', 'the highest leverage'])
plt.show()
```





```

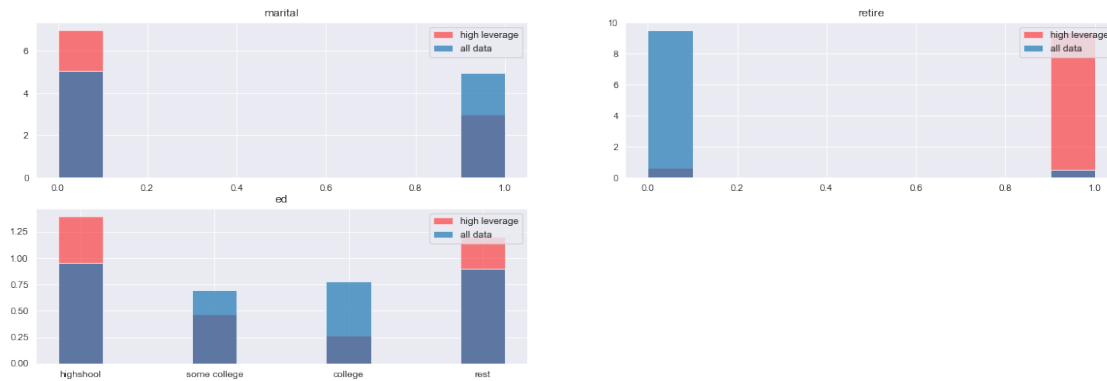
In [458]: temp = X.copy()
temp['HighLeverage'] = False
temp.loc[ind, 'HighLeverage'] = True
temp['ed'] = (1*temp.ed_highschool + 2*temp.ed_somecollege + 3*temp.ed_college + 3)

mpl.rcParams['figure.figsize'] = (20,10)
for i, r in enumerate(intervals_regressors):
    plt.subplot(len(intervals_regressors)//2+1, 2, i+1)
    plt.title(r)
    p1 = plt.violinplot(temp[temp['HighLeverage']==True][r],
                        showextrema=False, showmedians=True, vert=False)
    p2 = plt.violinplot(temp[r],
                        showextrema=True, showmedians=True, vert=False)
    for pc in p1['bodies']+p1['cmedians']:
        pc.set_facecolor('red')
        pc.set_edgecolor('red')
        pc.set_alpha(0.5)
    for pc in p2['bodies']+p2['cmedians']] + [p2['cbars']] + [p2['cmins']] + [p2['cmaxs']]:
        pc.set_facecolor('C0')
        pc.set_edgecolor('k')
        pc.set_alpha(0.7)
    p1['bodies'][0].set_label('high leverage')
    p2['bodies'][0].set_label('all data')
    plt.legend()
plt.show()

for i, r in enumerate([r for r in other_regressors + ['ed'] if r[:3]!='ed_']):
    plt.subplot(len(other_regressors)//2+1, 2, i+1)
    plt.title(r)
    plt.hist(temp[temp['HighLeverage']==True][r], alpha=0.5, density=True, color='r')
    plt.hist(temp[r], alpha=0.7, density=True)
    plt.legend(['high leverage', 'all data'], loc=1)
    if r=='ed':
        plt.xticks([0.15,1.05,1.95,2.85], ['highshool', 'some college', 'college', ''])
plt.show()

```





One can see from the distributions that the observations with the highest leverage have characteristics such as high tenure and long period at the current address. It is more often when the 'outlier' person is single, gets lower/no education and almost all of them are retired,

1.8 8.

Frequently data is missing. Pick up 5 rows in the data set and delete the value for address. Implement at least two approaches to fill in these values. Write down the corresponding formulas/model and give motivation for your approach. If you use standard routines then check how exactly the data imputation is implemented. How would you proceed if the value of the binary variable retire is missing? Implementation is not required.

In [480]: *# Missing completely at random*

```
X_miss = X.copy()
np.random.seed(100)
m = 5
ind_to_delete = np.random.randint(1, len(X_miss)+1, size=m)
X_miss.loc[ind_to_delete, 'address'] = np.nan
X_miss = get_normalized_X_y(X_miss, y, intervals_regressors)[0]
X_miss.loc[ind_to_delete, :]
```

```
Out[480]:
```

	tenure	retire	marital	ed_college	ed_highschool	ed_somcollege	\
521	0.490360	0	1	0	0	0	
793	0.256276	0	1	0	1	0	
836	-1.195048	0	1	0	1	0	
872	1.613966	0	0	0	0	0	
856	0.677628	0	0	0	0	0	

	address
521	NaN
793	NaN
836	NaN
872	NaN
856	NaN

1. Drop the rows with missing data:

```
In [481]: X1 = X_miss.drop(ind_to_delete, 0)
          y1 = y.drop(ind_to_delete, 0)

          LR1, beta_df1 = get_LR_beta(X1, y1, True)
```

```
LR scores:
R^2 = 0.715
R^2_adj = 0.711
AIC = -1.229
BIC = -1.160
```

2. Hot imputation: for each observation with missing address get the k=10 closest (by variables except address) observations without missing address and randomly choose one of them to replace missing value:

```
In [483]: np.random.seed(100)
          X2 = X_miss.copy()
          k = 10
          for i in range(m):
              other_values = list(X2.loc[ind_to_delete[i], :][other_regressors])
              X2[np.min(X2[other_regressors] - other_values == 0, 1)]

              intervals_regressors_ = sorted(list(set(intervals_regressors).difference({'address'})))
              intervals_values = list(X2.loc[ind_to_delete[i], :][intervals_regressors_])
              aux_df = (np.square(X2[intervals_regressors_] - intervals_values)).drop(ind_to_delete)
              aux_df['mean'] = np.mean(aux_df[intervals_regressors_], 1)
              ind_to_replace = np.random.choice(aux_df.sort_values(by='mean').head(k).index)
              # print(ind_to_delete[i], '<-', ind_to_replace)
              X2.loc[ind_to_delete[i], 'address'] = X2.loc[ind_to_replace, 'address']

          LR2, beta_df2 = get_LR_beta(X2, y, True)
```

```
LR scores:
R^2 = 0.715
R^2_adj = 0.711
AIC = -1.227
BIC = -1.158
```

3. Mean imputation: replace missing values with the mean of the variable:

```
In [484]: X3 = X_miss.copy()
          X3.loc[ind_to_delete, 'address'] = np.mean(X3['address'])
          LR3, beta_df3 = get_LR_beta(X3, y, True)
```

```
LR scores:
R^2 = 0.715
```

```
R^2_adj = 0.711
AIC = -1.227
BIC = -1.159
```

4. Regression imputation: replace missing values with the forecast using the Linear Regression model with address as a dependent variable, and all the others (except actual regressand log_longmon) as regressors:

```
In [485]: X4 = X_miss.copy()

XX = X_miss.drop(ind_to_delete, 0).drop('address', 1)
Xy = X_miss.drop(ind_to_delete, 0)['address']

XLR, _ = get_LR_beta(XX, Xy)
X4.loc[ind_to_delete, 'address'] = XLR.predict(X_miss.loc[ind_to_delete, :].drop('address', 1))
LR4, beta_df4 = get_LR_beta(X4, y, True)

LR scores:
R^2 = 0.715
R^2_adj = 0.711
AIC = -1.227
BIC = -1.159
```


All the methods seem to work well. Dropping 5 rows with missing values may not affect the result too much but in case of more frequent missing values we may lose useful information since it may be contained in the dropped rows. If the value of the binary variable retire is missing I would suggest to use hot imputation or median imputation or forecast using the logistic regression model.

1.9 9.

Now we look at the model assumptions. The variable address seems to be very significant. However, if we look at the residuals we observe that the variance of the residuals is rather different for different values of address. Run the Bartlett's test and compute the FGLS estimators assuming an exponential relationship between the variance of residuals and address. Compare the results with the original model. Explain the advantages of the (F)GLS estimation.

```
In [633]: represent_distribution(X['address'], varname='address')
q25, q75 = p_quantile(X.address, 0.25), p_quantile(X.address, 0.75)
upper_bound = q75 + 1.5*(q75 - q25)
print('We will drop observations with address value > {:.2f} (due to a small number of observations)')

Variable `address`
Number of observations: 1000
```



We will drop observations with address value > 2.87 (due to a small number of observations)


```
In [639]: mpl.rcParams['figure.figsize'] = (15,5)

temp = X[['address']]
temp['resid'] = LR_results.resid
temp = temp[temp.address < upper_bound]

vs = []
delta = 0.3
l = np.arange(-1,3,delta)
for i,j in zip(l, np.arange(0,4,delta)):
    v = temp[temp.address > i][temp.address < j].resid.var()
    vs.append(v)

fig, ax1 = plt.subplots()
ax1.scatter(temp.address, temp.resid, zorder=10)
ax1.set_ylabel('Residual')
ax1.set_xlabel('address')

ax2 = ax1.twinx()
ax2.grid('off')
ax2.plot(l,vs, 'r', label=' of residuals')
ax2.set_ylabel('Variance', color='r')
ax2.tick_params(axis='y', labelcolor='r')
plt.show()
```



```
In [657]: print('Bartlett\'s Test:')
          print('H0: sigma_1^2 = ... = sigma_2^2')
          print('H1: sigma_i^2 != sigma_j^2 for at least one pair (i,j)')
          stats.bartlett(*[r[1].values for r in temp.groupby('address').resid])
```

Bartlett's Test:

H0: $\sigma_1^2 = \dots = \sigma_2^2$

H1: $\sigma_i^2 \neq \sigma_j^2$ for at least one pair (i,j)

```
Out[657]: BartlettResult(statistic=38.69457145789063, pvalue=0.5289947719634349)
```

With the p-value is 0.529 we can not reject the null hypothesis. So we don't have enough arguments to reject that the variances are the same.

```
In [692]: multiplier = sm.OLS(np.log(temp.resid**2), sm.add_constant(temp.address)).fit().params
          print('The LR result: sigma_i^2 = sigma^2 * exp({:.3f} * address_i)\n'.format(multiplier))

          w = np.exp(multiplier * X.address)

          LR_OLS = sm.OLS(y, X_)
          LR_OLS_results = LR_OLS.fit()
          print(LR_OLS_results.summary())
          print()

          LR_FGLS = sm.GLS(y, X_, weights=w)
          LR_FGLS_results = LR_FGLS.fit()
          print(LR_FGLS_results.summary())
```

The LR result: $\sigma_i^2 = \sigma^2 * \exp(0.098 * \text{address}_i)$

OLS Regression Results

```
=====
Dep. Variable:          log_longmon    R-squared:          0.715
```

```

Model:                                OLS      Adj. R-squared:            0.713
Method:                             Least Squares  F-statistic:              355.7
Date:                               Tue, 26 Feb 2019  Prob (F-statistic):      2.36e-265
Time:                               17:25:40      Log-Likelihood:          -790.60
No. Observations:                    1000        AIC:                     1597.
Df Residuals:                        992         BIC:                     1636.
Df Model:                             7
Covariance Type:                     nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.1508	0.038	-3.969	0.000	-0.225	-0.076
tenure	0.8079	0.020	39.798	0.000	0.768	0.848
retire	0.2546	0.084	3.019	0.003	0.089	0.420
marital	0.0927	0.035	2.675	0.008	0.025	0.161
ed_college	0.1515	0.048	3.138	0.002	0.057	0.246
ed_highschool	0.1229	0.046	2.698	0.007	0.034	0.212
ed_somecollege	0.1061	0.050	2.136	0.033	0.009	0.204
address	0.0405	0.021	1.962	0.050	-4.99e-06	0.081
Omnibus:	2.592		Durbin-Watson:	1.944		
Prob(Omnibus):	0.274		Jarque-Bera (JB):	2.634		
Skew:	-0.054		Prob(JB):	0.268		
Kurtosis:	3.227		Cond. No.	6.24		

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

GLS Regression Results

```

Dep. Variable:          log_longmon      R-squared:            0.723
Model:                  GLS              Adj. R-squared:      0.721
Method:                 Least Squares    F-statistic:         370.2
Date:                   Tue, 26 Feb 2019  Prob (F-statistic):  1.66e-271
Time:                   17:25:40        Log-Likelihood:      -790.60
No. Observations:       1000            AIC:                1597.
Df Residuals:           992             BIC:                1636.
Df Model:                7
Covariance Type:        nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.1508	0.038	-3.969	0.000	-0.225	-0.076
tenure	0.8079	0.020	39.798	0.000	0.768	0.848
retire	0.2546	0.084	3.019	0.003	0.089	0.420
marital	0.0927	0.035	2.675	0.008	0.025	0.161
ed_college	0.1515	0.048	3.138	0.002	0.057	0.246

ed_highschool	0.1229	0.046	2.698	0.007	0.034	0.212
ed_somecollege	0.1061	0.050	2.136	0.033	0.009	0.204
address	0.0405	0.021	1.962	0.050	-4.99e-06	0.081
=====						
Omnibus:	2.592	Durbin-Watson:	1.944			
Prob(Omnibus):	0.274	Jarque-Bera (JB):	2.634			
Skew:	-0.054	Prob(JB):	0.268			
Kurtosis:	3.227	Cond. No.	6.24			
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The advantage of the feasible GLS is that the parameter estimator is more efficient than that of OLS. This advantage is seen only when the weights of the error covariance matrix are correlated with data. If the weights are known, we have the best linear unbiased estimator. In our case and more often it is unknown. Still with the estimated weights (FGLS) we may get more efficient parameters.

Here we got identical results which is expectable since with the Bartlett's test we couldn't reject the hypothesis that the variances are the same.

1.10 10.

Compute the White estimator of covariance matrix of the OLS estimators. Run the t-tests and compare the results with the original model. Explain the advantages of the White estimator for the variance.

```
In [705]: print(LR_OLS.fit().summary())
          print()
          print(LR_OLS.fit(cov_type='HCO').summary())
```

OLS Regression Results						
=====						
Dep. Variable:	log_longmon	R-squared:	0.715			
Model:	OLS	Adj. R-squared:	0.713			
Method:	Least Squares	F-statistic:	355.7			
Date:	Tue, 26 Feb 2019	Prob (F-statistic):	2.36e-265			
Time:	17:32:12	Log-Likelihood:	-790.60			
No. Observations:	1000	AIC:	1597.			
Df Residuals:	992	BIC:	1636.			
Df Model:	7					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-0.1508	0.038	-3.969	0.000	-0.225	-0.076
tenure	0.8079	0.020	39.798	0.000	0.768	0.848

retire	0.2546	0.084	3.019	0.003	0.089	0.420
marital	0.0927	0.035	2.675	0.008	0.025	0.161
ed_college	0.1515	0.048	3.138	0.002	0.057	0.246
ed_highschool	0.1229	0.046	2.698	0.007	0.034	0.212
ed_somecollege	0.1061	0.050	2.136	0.033	0.009	0.204
address	0.0405	0.021	1.962	0.050	-4.99e-06	0.081

```

=====
Omnibus:                2.592    Durbin-Watson:                1.944
Prob(Omnibus):          0.274    Jarque-Bera (JB):        2.634
Skew:                   -0.054    Prob(JB):                0.268
Kurtosis:               3.227    Cond. No.                6.24
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```

=====
Dep. Variable:          log_longmon    R-squared:                0.715
Model:                  OLS            Adj. R-squared:           0.713
Method:                 Least Squares   F-statistic:             324.6
Date:                  Tue, 26 Feb 2019 Prob (F-statistic):       1.88e-251
Time:                  17:32:12         Log-Likelihood:          -790.60
No. Observations:      1000            AIC:                    1597.
Df Residuals:          992             BIC:                    1636.
Df Model:               7
Covariance Type:       HC0
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-0.1508	0.041	-3.649	0.000	-0.232	-0.070
tenure	0.8079	0.021	37.769	0.000	0.766	0.850
retire	0.2546	0.087	2.916	0.004	0.084	0.426
marital	0.0927	0.035	2.644	0.008	0.024	0.161
ed_college	0.1515	0.048	3.129	0.002	0.057	0.246
ed_highschool	0.1229	0.047	2.625	0.009	0.031	0.215
ed_somecollege	0.1061	0.051	2.073	0.038	0.006	0.206
address	0.0405	0.021	1.934	0.053	-0.001	0.081

```

=====
Omnibus:                2.592    Durbin-Watson:                1.944
Prob(Omnibus):          0.274    Jarque-Bera (JB):        2.634
Skew:                   -0.054    Prob(JB):                0.268
Kurtosis:               3.227    Cond. No.                6.24
=====

```

Warnings:

[1] Standard Errors are heteroscedasticity robust (HC0)

The white estimator corrected the confidence intervals since it is computed based on the variance of the parameters — white estimator of the covariance matrix based on the residuals and is used in computing the variance of the parameters. Also in the significance tests the statistic values changed and the p-values increased, but not significantly.

Problem 5: further issues

(a) Let $y_i = \beta_0 + \sum_{k=1}^K \beta_k x_{ik} + u_i$, $i = \overline{1, N}$ – a linear regression model for y_i .

$$\hat{y}_i = \hat{\beta}_0 + \sum_{k=1}^K \beta_k x_{ik} \quad \hat{u}_i = y_i - \hat{y}_i$$

Let $y_i^* = y_i + 10$. To get the same least squares of residuals:

$$\hat{u}_i^* = \hat{u}_i = (y_i + 10 - 10) - \hat{y}_i = (y_i + 10) - (\hat{y}_i + 10) = y_i^* - \hat{y}_i^*$$

we should have $\hat{y}_i^* = \hat{y}_i + 10 = (\hat{\beta}_0 + 10) + \sum_{k=1}^K \beta_k x_{ik}$. So $\hat{\beta}_0^* = \hat{\beta}_0 + 10$

$$\bar{y}^* = \frac{\sum_{i=1}^N y_i^*}{N} = \bar{y} + 10 \quad s_{y^*}^2 = \frac{\sum_{i=1}^N (y_i^* - \bar{y}^*)^2}{N-1} = \frac{\sum_{i=1}^N (y_i - \bar{y})^2}{N-1} = s_y^2$$

$$\Rightarrow R^{*2} = 1 - \frac{\sum_{i=1}^N \hat{u}_i^{*2}}{s_{y^*}^2} = 1 - \frac{\sum_{i=1}^N \hat{u}_i^2}{s_y^2} = R^2$$

Let's also show it in matrix form:

$$\mathbf{y} = X\boldsymbol{\beta} + \mathbf{u}, \quad \hat{\mathbf{y}} = X\hat{\boldsymbol{\beta}}, \quad \hat{\mathbf{u}} = \mathbf{y} - \hat{\mathbf{y}}$$

$$X \in N \times (K+1), \quad \mathbf{y}, \hat{\mathbf{y}}, \mathbf{u} \in \mathbb{R}^N, \quad \boldsymbol{\beta}, \hat{\boldsymbol{\beta}} \in \mathbb{R}^{K+1}$$

The paramaters from OLS:

$$\hat{\boldsymbol{\beta}} = (X^\top X)^{-1} X^\top \mathbf{y}$$

Let $(X^\top X)^{-1} = \mathbf{A} = \begin{pmatrix} \mathbf{a}_0^\top \\ \vdots \\ \mathbf{a}_K^\top \end{pmatrix}$. Then $\hat{\boldsymbol{\beta}} = \mathbf{A}X^\top \mathbf{y}$, $\hat{\beta}_0 = \mathbf{a}_0^\top X^\top \mathbf{y}$.

Let $y_i^* = y_i + 10$.

$$\mathbf{y}^* = \mathbf{y} + 10 \cdot \mathbf{1}, \quad \mathbf{1} = \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix} \in \mathbb{R}^N$$

$$\hat{\boldsymbol{\beta}}^* = \mathbf{A}X^\top \mathbf{y}^* = \mathbf{A}X^\top \mathbf{y} + 10\mathbf{A}X^\top \mathbf{1} \quad \hat{\beta}_0^* = \hat{\beta}_0 + 10 \mathbf{a}_0^\top X^\top \mathbf{1}$$

$$\hat{\mathbf{y}}^* = X\hat{\boldsymbol{\beta}}^* = X\hat{\boldsymbol{\beta}} + 10X\mathbf{A}X^\top \mathbf{1} = \hat{\mathbf{y}} + 10X\mathbf{A}X^\top \mathbf{1}$$

Let's prove that $X\mathbf{A}X^\top \mathbf{1} = \mathbf{1}$. Indeed, since $P = X(X^\top X)^{-1}X^\top$ is the projector onto the column space $\mathcal{C}(X)$ and the first column is $\mathbf{1}$: $P\mathbf{1} = \mathbf{1}$. So:

$$\hat{\mathbf{y}}^* = \hat{\mathbf{y}} + 10 \cdot \mathbf{1}$$

From this one can also see that $\mathbf{A}X^\top \mathbf{1} = \mathbf{e}_1 = (1 \ 0 \ \cdots \ 0)^\top \in \mathbb{R}^{K+1}$ as the coefficients for the linear combination of columns of X to get $\mathbf{1}$, and hence:

$$\hat{\boldsymbol{\beta}}^* = \mathbf{A}X^\top \mathbf{y} + 10\mathbf{e}_1 \quad \hat{\beta}_0^* = \hat{\beta}_0 + 10$$

$$\hat{\mathbf{u}}^* = \mathbf{y}^* - \hat{\mathbf{y}}^* = \mathbf{y} + 10 \cdot \mathbf{1} - \hat{\mathbf{y}} - 10 \cdot \mathbf{1} = \hat{\mathbf{u}}$$

$$R^{*2} = 1 - \frac{\|\hat{\mathbf{u}}^*\|^2}{s_{y^*}^2} = 1 - \frac{\|\hat{\mathbf{u}}\|^2}{s_y^2} = R^2$$

Answer: $\hat{\beta}_0^* - \hat{\beta}_0 = 10$;

$$R^{*2} = R^2;$$

Now let $y_i^* = y_i - 10$. Similarly $\hat{\boldsymbol{\beta}}^* = \hat{\boldsymbol{\beta}} - 10\mathbf{e}_1$,

$$\hat{\beta}_0^* - \hat{\beta}_0 = -10 \quad R^{*2} = R^2$$

So for δ_y we have $y_i^* = y_i + \delta_y$: $\hat{\beta}_0^* = \hat{\beta}_0 + \delta_y$, $\hat{\beta}_i^* = \hat{\beta}_i$ ($i = \overline{1, K}$), $R^{*2} = R^2$

(b) Now let's continue with a linear regression model:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \cdots + \hat{\beta}_K x_{iK}$$

If we change say all x_{i1} by δ_1 : $x_{i1}^* = x_{i1} + \delta_1$, similarly from the OLS we should get the same $\hat{y}_i^* = \hat{y}_i$, since the change of x_{i1} is linear. So:

$$\begin{aligned} \hat{y}_i &= \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \cdots + \hat{\beta}_K x_{iK} = \hat{\beta}_0 + \hat{\beta}_1 (x_{i1} + \delta_1 - \delta_1) + \cdots + \hat{\beta}_K x_{iK} = \hat{\beta}_0 + \hat{\beta}_1 x_{i1}^* - \hat{\beta}_1 \delta_1 + \cdots + \hat{\beta}_K x_{iK} = \\ &= (\hat{\beta}_0 - \hat{\beta}_1 \delta_1) + \hat{\beta}_1 x_{i1}^* + \cdots + \hat{\beta}_K x_{iK} = \hat{\beta}_0^* + \hat{\beta}_1^* x_{i1}^* + \cdots + \hat{\beta}_K^* x_{iK} \\ &\hat{\beta}_0^* = \hat{\beta}_0 - \hat{\beta}_1 \delta_1 \end{aligned}$$

So in general for $x_{ik}^* = x_{ik} + \delta_k$:

$$\begin{aligned} \hat{\beta}_0^* &= \hat{\beta}_0 - \sum_{k=1}^K \hat{\beta}_k \delta_k \\ \hat{u}_i^* &= \hat{u}_i \Rightarrow R^{*2} = R^2 \end{aligned}$$

(c) Using the above conclusions, we can get the result of demeaning: $y_i^* = y_i - \bar{y}$, $x_{ik}^* = x_{ik} - \bar{x}_k$, $k = \overline{1, K}$, $i = \overline{1, N}$. Now we will have:

$$\hat{\beta}_0^* = \hat{\beta}_0 + \sum_{k=1}^K \hat{\beta}_k \bar{x}_k - \bar{y} \quad R^{*2} = R^2$$